

Take Home Test #3
Dot Product Optimization
Anthony Ramos
CSC342/43
May 16th, 2021
Professor Isidor Gertner

Contents

Objective	3
Introduction.....	3
<i>My CPU Info</i>	<i>3</i>
<i>Dot Product Program</i>	<i>5</i>
<i>Test File</i>	<i>5</i>
Dot Product (Array Indexing) (Not Optimized).....	6
Dot Product (Array Pointer Arithmetic) (Not Optimized)	7
Enabling Compiler Optimizations.....	8
Dot Product (Array Indexing) (Compiler Optimized).....	9
Dot Product (Array Pointer Arithmetic) (Compiler Optimized)	10
Comparing Dot Product Assembler Source	11
Dot Product (Vector Instructions) (Pointer)	14
Optimizing Dot Product in A Linux Environment	16
<i>Enabling Compiler Optimization Flags.....</i>	<i>16</i>
<i>Test File</i>	<i>17</i>
<i>Comparing Assembly Source Files</i>	<i>19</i>
<i>Results</i>	<i>22</i>

Objective

The objective of this assignment is to implement a program to compute the dot product using array indexing and pointer arithmetic. Next, optimization features will be applied (by software) followed by a manually optimized approach on the compiler generated assembly code. In all three cases, the execution times will be measured and analyzed. The simulation environments utilized were Intel X86 32-Bit compiler in MS Visual Studio and Intel X86 64-Bit GCC compiler in a Linux (18.04.5) environment.

Introduction

My CPU Info

Prior writing and analyzing any code, it was of interest to learn about the CPU capabilities in my hardware. To identify my hardware's processor, I utilized the *CPUID* application. Figure 1a shows information regarding my processor along with the instruction sets that are available (in red) for executing.

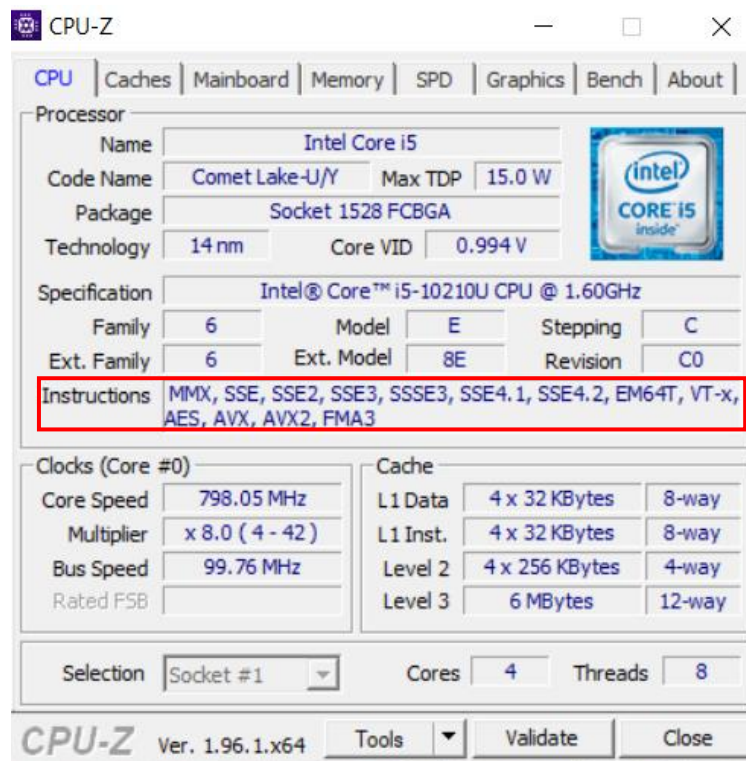


Figure 1a: My CPU Processor Info from CPUID Application

To confirm that my processor was properly detected, I checked for it in Linux. It can be confirmed that my processor is a 10th Generation Intel Core i5.

```
anthony@compOrg:~/THT3$ lscpu | grep "Model name"
Model name:      Intel(R) Core(TM) i5-10210U CPU @ 1.60GHz
```

Figure 1b: My CPU Processor Info from CPUID Application

To determine my processor's capabilities, I did some research online and found [data](#) regarding my processor's performance in several categories as well as its performance relative to other popular processors.

CPU Test Suite Average Results for Intel Core i5-10210U @ 1.60GHz	
Integer Math	23,993 MOps/Sec
Floating Point Math	14,603 MOps/Sec
Find Prime Numbers	18 Million Primes/Sec
Random String Sorting	12 Thousand Strings/Sec
Data Encryption	2,176 MBytes/Sec
Data Compression	84.7 MBytes/Sec
Physics	424 Frames/Sec
Extended Instructions	5,513 Million Matrices/Sec
Single Thread	2,268 MOps/Sec

From submitted results to PerformanceTest V10 as of 13th of May 2021.

Figure 1c: My CPU Processor Capabilities




Figure 1d: My CPU Processor Performance among Popular Processors

Dot Product Program

Two programs to compute the dot product will be written using two programming structures. One program will implement the dot product by utilizing array index notation as shown in figure 2a. The second will utilize pointer arithmetic as shown in figure 2b.

```
int DotProduct_Index(int a1[], int a2[], int n) {
    int sum = 0;
    for (int i = 0; i < n; i++)
        sum += (a1[i] * a2[i]);
    return sum;
}
```

Figure 2a: Dot Product Program (Array Index Notation)

```
int DotProduct_Pointer(int *a1, int *a2, int n) {
    int *a, *b;
    int sum = 0;
    for (a = &a1[0], b = &a2[0]; a < &a1[n]; a++, b++)
        sum += ((*a) * (*b));
    return sum;
}
```

Figure 2b: Dot Product Program (Pointer Arithmetic)

Test File

A third program (test file) will be utilized to serve as the *main* program that will call the previous two programs (one per run). It should be noted we care not for the actual result of the dot product. What matters is the performance (i.e., execution time of each program). To measure the execution time, we will utilize a highly accurate timing mechanism called the `QueryPerformanceCounter()` function.

```
if (QueryPerformanceCounter((LARGE_INTEGER*)&timeStart) != 0) {
    result = DotProduct_Index(arr1, arr2, arraySize);
    //result = DotProduct_Pointer(&arr1[0], &arr2[0], arraySize);
    QueryPerformanceCounter((LARGE_INTEGER*)&timeEnd);

    QueryPerformanceFrequency((LARGE_INTEGER*)&freq);

    cout << "QueryPerformance frequency: " << freq << " counts per second" <<
endl;

    cout << "Array size: " << arraySize << endl;
    cout << "Elapsed Time: " << (timeEnd - timeStart) * 1.0 / freq << "
seconds" << endl;
}
else {
    DWORD dwError = GetLastError();
    cout << "Error = " << dwError << endl;
}
```

Figure 3: `QueryPerformanceCounter()`

Dot Product (Array Indexing) (Not Optimized)

In the first round of simulations, we disable some optimization features (*Automatic Parallelization & Automatic Vectorization*). Utilizing the `QueryPerformanceCounter()`, several measurements were taken to evaluate the program's performance. The average execution times of multiple runs for array sizes of $n = 2^4$ to 2^{20} will then be plotted. Figure 5 shows the plotted results are linear which is expected.

Array Size	Execution Time
16	0.00000036
32	0.00000041
64	0.00000054
128	6.71429E-07
256	0.00000124
512	0.00000198
1024	0.00000302
2048	0.00000616
4096	0.00001308
8192	0.00002178
16384	3.93556E-05
32768	0.00008305
65536	0.0001780
131072	0.0002889
262144	0.0007395
524288	0.001501317
1048576	0.00270965

Table 1: Execution times of dot product program via array indexing (Not-Optimized)




Figure 4: Execution Times of Dot Product Via Array Indexing (No Optimization)

Dot Product (Array Pointer Arithmetic) (Not Optimized)

The previous process is repeated for the second program using pointer arithmetic. The execution times are shown in table 2 and plotted results in figure 6. Again, the expected plot is linear.

Array Size	Execution Time
16	0.00000034
32	0.0000004
64	0.00000052
128	0.00000072
256	0.00000096
512	2.1660E-06
1024	3.6625E-06
2048	6.1375E-06
4096	1.38333E-05
8192	2.90111E-05
16384	0.00004735
32768	9.54111E-05
65536	1.92E-04
131072	4.52E-04
262144	8.12E-04
524288	1.65E-03
1048576	3.31E-03

Table 2: Execution times of dot product program via array pointer arithmetic (Not-Optimized)




Figure 6: Execution Times of Dot Product Via Array Indexing (No Optimization)

Enabling Compiler Optimizations

Next, we wish to improve the execution time of each program by enabling the aforementioned optimization settings within Visual Studio. Further information regarding these settings can be found [here](#).

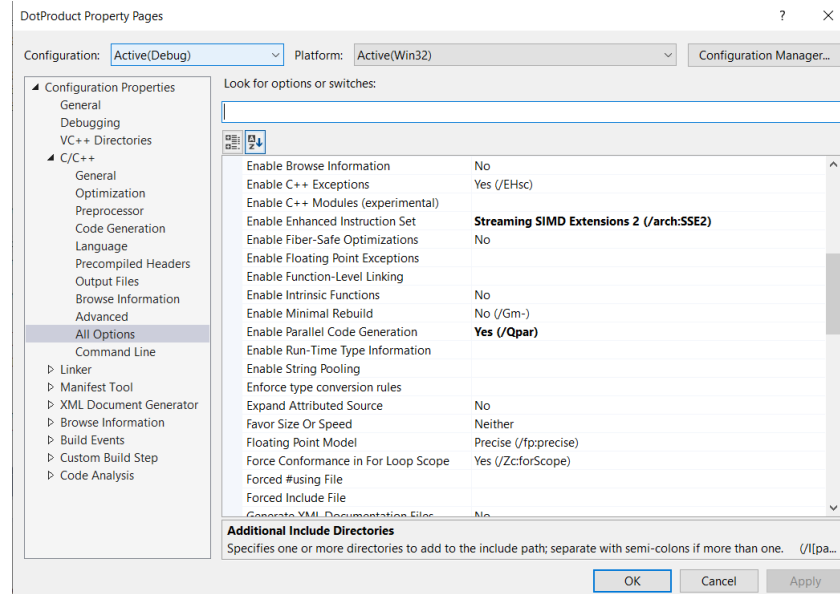


Figure 7: Enabling Optimization Features in Visual Studio

Dot Product (Array Indexing) (Compiler Optimized)

With optimization settings set, we repeat the same measurements for array sizes $n = 2^4$ to 2^{20} and tabulate the results as shown in table 3. Figure 8 shows the plotted results.

Array Size	Execution Time
16	0.000000325
32	0.00000039
64	0.00000052
128	0.00000065
256	0.00000111
512	0.00000167
1024	0.00000297
2048	0.00000502
4096	0.0000126
8192	0.00001731
16384	0.00003764
32768	0.00007569
65536	0.000171963
131072	0.00024365
262144	0.00065296
524288	0.001303629
1048576	0.002459667

Table 3: Execution Times of Dot Product Program Via Array Indexing (Compiler-Optimized)




Figure 8: Execution Times of Dot Product Via Array Indexing (Compiler-Optimized)

Dot Product (Array Pointer Arithmetic) (Compiler Optimized)

Using the same optimization settings, we repeat the previous process and measure the execution times as described in table 4. The results are plotted in figure 9.

Array Size	Execution Time
16	0.0000003
32	0.00000042
64	0.00000042
128	0.00000066
256	0.00000096
512	0.00000153
1024	0.000002825
2048	0.000005825
4096	0.000009625
8192	2.17375E-05
16384	0.000038475
32768	8.09875E-05
65536	0.00018367
131072	0.0004259
262144	0.000789625
524288	0.001547517
1048576	0.00245714

Table 4: Execution Times of Dot Product Program Via Pointer Arithmetic (Compiler-Optimized)




Figure 9: Execution Times of Dot Product Via Pointer Arithmetic (Compiler-Optimized)

Comparing Dot Product Assembler Source

The application of optimization settings often leads to a reduction in execution time (in our case). While the high-level programming remains unchanged, the assembly source is altered to yield better performance. Figure 10a shows the partial assembly code for the dot product program via array indexing for when no optimization (left) and optimization (right) is applied. It can be seen that, with optimization settings, the assembly source becomes partially vectorized. Although the assembly becomes longer and can be more difficult to debug, it greatly improves the programs performance by reducing execution time.

<pre> int sum = 0; 00B11743 mov dword ptr [sum],0 #pragma loop(hint_parallel(8)) for (int i = 0; i < n; i++) 00B1174A mov dword ptr [ebp-8],0 00B11751 jmp DotProduct_Index+2Ch (0B1175Ch) 00B11753 mov eax,dword ptr [ebp-8] 00B11756 add eax,1 00B11759 mov dword ptr [ebp-8],eax 00B1175C mov eax,dword ptr [ebp-8] 00B1175F cmp eax,dword ptr [n] 00B11762 jge DotProduct_Index+4Fh (0B1177Fh) sum += (a1[i] * a2[i]); 00B11764 mov eax,dword ptr [ebp-8] 00B11767 mov ecx,dword ptr [a1] 00B1176A mov edx,dword ptr [ebp-8] 00B1176D mov esi,dword ptr [a2] 00B11770 mov eax,dword ptr [ecx+eax*4] 00B11773 imul eax,dword ptr [esi+edx*4] 00B11777 add eax,dword ptr [sum] 00B1177A mov dword ptr [sum],eax 00B1177D jmp DotProduct_Index+23h (0B11753h) return sum; 00B1177F mov eax,dword ptr [sum] } </pre>	<pre> #pragma loop(hint_parallel(8)) for (int i = 0; i < n; i++) 00221755 cmp edi,8 00221758 jnb DotProduct_Index+0B6h (02217E6h) 0022175E cmp dword ptr [__isa_available (0BC2AD4h)],2 00221765 jnl DotProduct_Index+0B6h (02217E6h) 00221768 mov ecx,edi 0022176D and ecx,80000007h 00221773 jns DotProduct_Index+4Ah (022177Ah) 00221775 dec ecx 00221776 or ecx,0FFFFFFFh 00221779 inc ecx 0022177A mov edx,dword ptr [a1] 0022177D mov esi,edi 0022177F mov ebx,dword ptr [a2] 00221782 sub esi,ecx 00221784 mov dword ptr [ebp-4],ebx 00221787 xorps xmm3,xmm3 0022178A sub dword ptr [ebp-4],edx 0022178D xorps xmm2,xmm2 00221790 lea ecx,[edx+10h] 00221793 mov edx,dword ptr [ebp-4] 00221796 movups xmm0,xmmword ptr [ebx+eax*4] 0022179A add eax,8 0022179D lea ecx,[ecx+20h] sum += (a1[i] * a2[i]); 002217A0 movups xmm1,xmmword ptr [ecx-30h] 002217A4 pmulld xmm1,xmm0 002217A9 padd xmm3,xmm1 002217AD movups xmm0,xmmword ptr [edx+ecx-20h] 002217B2 movups xmm1,xmmword ptr [ecx-20h] 002217B6 pmulld xmm1,xmm0 002217BB padd xmm2,xmm1 002217BF cmp eax,esi 002217C1 jnl DotProduct_Index+66h (0221796h) 002217C3 padd xmm2,xmm3 002217C7 movaps xmm0,xmm2 002217CA psrldq xmm0,8 002217CF padd xmm2,xmm0 002217D3 movups xmm0,xmm2 002217D6 psrldq xmm0,4 002217DB padd xmm2,xmm0 002217DF movd esi,xmm2 002217E3 mov dword ptr [sum],esi </pre>
--	---

Figure 10a: Assembly Source for Dot Product via Array Indexing (Partial)

A similar observation can be seen for the dot product via pointer arithmetic. Although, the compiler optimized assembly source did not utilize vector instructions. Instead, it optimized the assembly code by reducing the number of machine instructions required to perform the same task.

```

    int *a, *b;
    int sum = 0;
008317C3  mov     dword ptr [sum],0
    #pragma loop(hint_parallel(8))
    for (a = &a1[0], b = &a2[0]; a < &a1[n]; a++, b++)
008317CA  mov     eax,4
008317CF  imul    ecx,eax,0
008317D2  add     ecx,dword ptr [a1]
008317D5  mov     dword ptr [a],ecx
008317D8  mov     edx,4
008317DD  imul    eax,edx,0
008317E0  add     eax,dword ptr [a2]
008317E3  mov     dword ptr [b],eax
008317E6  jmp     DotProduct_Pointer+4Ah (08317FAh)
008317E8  mov     eax,dword ptr [a]
008317EB  add     eax,4
008317EE  mov     dword ptr [a],eax
008317F1  mov     ecx,dword ptr [b]
008317F4  add     ecx,4
008317F7  mov     dword ptr [b],ecx
008317FA  mov     eax,dword ptr [n]
008317FD  mov     ecx,dword ptr [a1]
00831800  lea     edx,[ecx+eax*4]
00831803  cmp     dword ptr [a],edx
00831806  jae     DotProduct_Pointer+6Bh (083181Bh)

    sum += ((*a) * (*b));
00831808  mov     eax,dword ptr [a]
0083180B  mov     ecx,dword ptr [b]
0083180E  mov     edx,dword ptr [eax]
00831810  imul    edx,dword ptr [ecx]
00831813  add     edx,dword ptr [sum]
00831816  mov     dword ptr [sum],edx
00831819  jmp     DotProduct_Pointer+38h (08317E8h)
    return sum;
0083181B  mov     eax,dword ptr [sum]
}

    int *a, *b;
    int sum = 0;
    #pragma loop(hint_parallel(8))
    for (a = &a1[0], b = &a2[0]; a < &a1[n]; a++, b++)
002A18E9  and     eax,3FFFFFFFh
002A18EE  cmp     eax,2
002A18F1  jnb     DotProduct_Pointer+66h (02A1926h)
002A18F3  mov     eax,esi
002A18F5  sub     eax,ecx
002A18F7  mov     dword ptr [a1],eax
002A18FA  mov     edx,eax
002A18FC  nop     dword ptr [eax]
    sum += ((*a) * (*b));
002A1900  mov     eax,dword ptr [esi]
002A1902  add     esi,8
002A1905  imul    eax,dword ptr [ecx]
002A1908  add     edi,eax
002A190A  mov     eax,dword ptr [edx+ecx+4]
002A190E  imul    eax,dword ptr [ecx+4]
002A1912  add     ecx,8
002A1915  mov     edx,dword ptr [ebp+8]
002A1918  add     ebx,eax
002A191A  mov     eax,dword ptr [ebp+10h]
002A191D  add     eax,0FFFFFFFh
002A1920  cmp     ecx,eax
002A1922  jnb     DotProduct_Pointer+40h (02A1900h)
002A1924  xor     edx,edx

```

Figure 10b: Assembly Source for Dot Product via Pointer Arithmetic (Partial)

Figures 11a and 11b compare the execution times for both dot product programs. The results confirm our earlier analysis. Optimization settings did reduce the execution time for both cases. This of course, was most noticeable for very large array sizes.



Figure 11a: Comparison of Execution Times for Dot Product via Array Indexing




Figure 11b: Comparison of Execution Times for Dot Product via Pointer Arithmetic

Dot Product (Vector Instructions) (Pointer)

In measuring the execution times after applying built in optimization settings, it was now time to attempt to optimize the compiler generated code ourselves. However, difficulty arose, and I was not able to successfully optimize the existing code manually. Instead, I moved towards utilizing SSE3 vector instructions from previous Intel tutorial to implement the dot product from scratch whose code is shown below.

```
float DotProduct_Vector(float *a1, float *a2, int n) {
    float x = 0.0;
    _asm {
        pxor xmm0, xmm0; initialize xmm0 to 0 (xmm0 stores dot product result)

        mov eax, dword ptr[a1]; %eax points to arr1[]
        mov ebx, dword ptr[a2]; %ebx points to arr2[]
        mov ecx, dword ptr[n]; Array Size

        $MyLoop:
        movups xmm0, [eax]
        movups xmm1, [ebx]
        mulps  xmm1, xmm2; Compute arr1[i] * arr2[i]
        addps  xmm0, xmm1; Compute x + arr1[i]* arr2[i]

        add eax, 16
        add ebx, 16
        sub ecx, 4; Loop - 4
        jnz $MyLoop ; Loop if ecx is not zero

        haddps xmm0, xmm0 ; Horizontal Add
        haddps xmm0, xmm0 ; Horizontal Add
        movss dword ptr[x], xmm0 ; Store result in x
    }
    return 0;
}
```

Figure 12: Implementation of Dot Product Using Vector Instructions

We perform our previous measurements which yield the results in table 5.

Array Size	Execution Time
16	4.42857E-07
32	5.14286E-07
64	5.42857E-07
128	0.0000006
256	6.71429E-07
512	0.0000007
1024	1.06667E-06
2048	1.25556E-06
4096	0.0000022
8192	0.0000048
16384	8.27778E-06
32768	0.0000192
65536	3.98889E-05
131072	0.000105513
262144	0.000184338
524288	0.000427225
1048576	0.000837667

Table 5: Execution Times of Dot Product Program Via Pointer Arithmetic (SSE3 Vector Instructions)

Finally, the plot in figure 13 compares the previous plots with that obtained from the data in table 5. When applying the automatic parallelization and vectorization settings, the execution time is greatly reduced. This is most noticeable for larger inputs. The use of SSE3 vector instructions provided a significant advantage in reducing execution time. However, a similar implementation for the dot product via array indexing could not be realized.




Figure 13: Summary of Dot Product (via Pointer Arithmetic) Execution Times

Optimizing Dot Product in A Linux Environment

We now move towards optimization in a Linux Environment. While one may attempt to manually optimize the compiler generated assembly code produced, GCC has readily available optimization features via *compiler flags*.

Enabling Compiler Optimization Flags

There are many optimization flags of which a complete listing can be found here. However, we will only consider the `-O` optimization flag. Specifically, we will consider the flags that reduce execution time of which there are three (and a no optimization default) and described in table 5 below. The '+' symbol indicates an increase an execution time while the '-' symbol indicates a decrease in execution time. The '--' and '---' indicate even further decreases in execution time.

Optimization Flag	Optimization Level	Execution Time	Notes
<code>-O0</code>	0 (Default)	+	No Optimization
<code>-O1</code>	1	-	Basic Optimization
<code>-O2</code>	2	--	Recommend Level
<code>-O3</code>	3 (Highest)	---	Optimization Not Guaranteed

Table 5: Summary of Optimization Flags

To set the optimization flag, we utilize the following syntax:

```
gcc -Olevel [source files] -o [output filename]
```

There are tradeoffs to applying these optimization flags – most notably in compilation time and memory usage. An increase in optimization leads to an increase in both these metrics. Utilizing higher optimization levels may not guarantee better performance and may even give incorrect or unreliable results according to sources found online. Nevertheless, we shall apply all four optimization levels and analyze the execution times.

```
anthony@comp0rg:~/THT3$ g++ -O0 TestFile.cpp DotProduct_Index.cpp DotProduct_Pointer.cpp -o performanceTest
anthony@comp0rg:~/THT3$ g++ -O1 TestFile.cpp DotProduct_Index.cpp DotProduct_Pointer.cpp -o performanceTestO1
anthony@comp0rg:~/THT3$ g++ -O2 TestFile.cpp DotProduct_Index.cpp DotProduct_Pointer.cpp -o performanceTestO2
anthony@comp0rg:~/THT3$ g++ -O3 TestFile.cpp DotProduct_Index.cpp DotProduct_Pointer.cpp -o performanceTestO3
```

Figure 14a: Compiling Dot Product Programs Utilizing Optimization Flags

Test File

Figure 14b shows the test file utilized to measure the execution times of each dot product program. Specifically, for a given array size, 1000 measurements will be taken and the average execution time will be outputted.

```
#include <iostream>
#include <time.h>
#include <stdint.h>
#include <stdlib.h>

extern int DotProduct_Index(int a1[], int a2[], int n);
extern int DotProduct_Pointer(int *a1, int *a2, int n);

#define NANO 1000000000L
using namespace std;

const int runs = 1000;
const int arraySize = 256;
static int arr1[arraySize];
static int arr2[arraySize];

int main() {
    uint64_t elapsed_time;
    uint64_t PointerSum = 0;
    uint64_t IndexSum = 0;
    int result;
    struct timespec timeStart, timeEnd;

    for (int i = 0; i < arraySize; i++) {
        arr1[i] = i;
        arr2[i] = i/2;
    }

    cout << "Array Size = " << arraySize << endl;

    for (int i = 0; i < runs; i++) {
        clock_gettime(CLOCK_PROCESS_CPUTIME_ID, &timeStart);
// Begin Timer
        result = DotProduct_Pointer(&arr1[0], &arr2[0],
arraySize);
        clock_gettime(CLOCK_PROCESS_CPUTIME_ID, &timeEnd); //
End Timer

        elapsed_time = NANO * (timeEnd.tv_sec -
timeStart.tv_sec) + (timeEnd.tv_nsec - timeStart.tv_nsec);
        //cout << "Execution Time = " << elapsed_time << "ns"
<< endl;

        PointerSum += elapsed_time;
        //cout << PointerSum << endl;
    }
}
```

```

cout << "Dot Product Pointer :: Average execution time after " << runs << "
runs = " << PointerSum/runs << " ns" << endl;

for (int j = 0; j < runs; j++) {
    clock_gettime(CLOCK_PROCESS_CPUTIME_ID, &timeStart); // Begin
Timer
    result = DotProduct_Index(arr1, arr2, arraySize);
    clock_gettime(CLOCK_PROCESS_CPUTIME_ID, &timeEnd); // End
Timer

    elapsed_time = NANO * (timeEnd.tv_sec - timeStart.tv_sec) +
(timeEnd.tv_nsec - timeStart.tv_nsec);
    //cout << "Execution Time = " << elapsed_time << "ns" <<
endl;

    IndexSum += elapsed_time;
    //cout << IndexSum << endl;

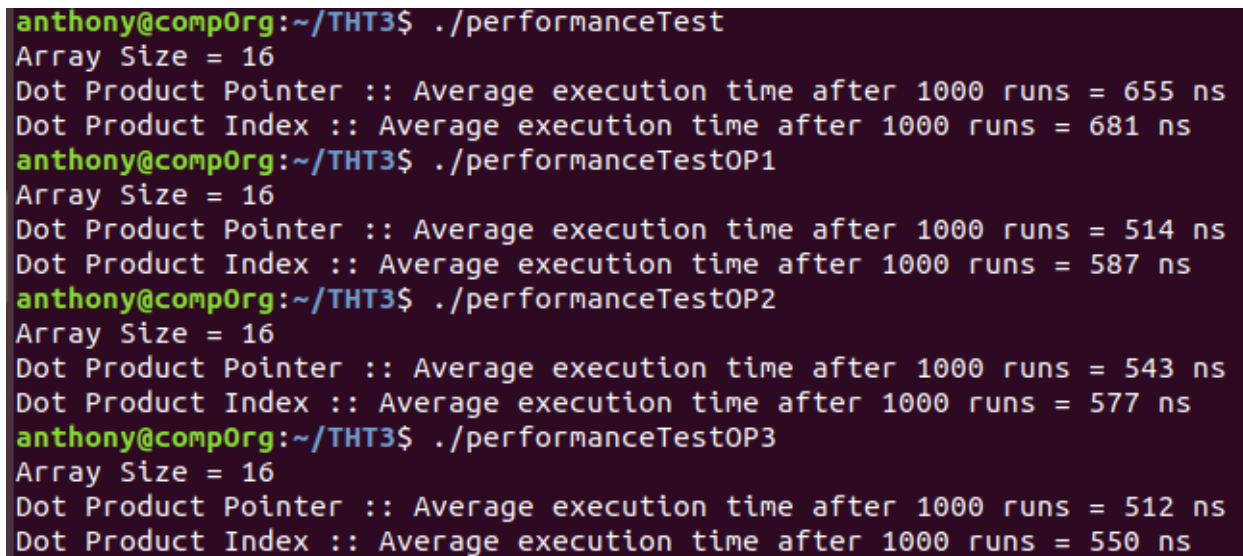
}

cout << "Dot Product Index :: Average execution time after " << runs
<< " runs = " << IndexSum/runs << " ns" << endl;
return 0;

}

```

Figure 14b: Test Program to Measure Average Execution Times



```

anthony@compOrg:~/THT3$ ./performanceTest
Array Size = 16
Dot Product Pointer :: Average execution time after 1000 runs = 655 ns
Dot Product Index :: Average execution time after 1000 runs = 681 ns
anthony@compOrg:~/THT3$ ./performanceTestOP1
Array Size = 16
Dot Product Pointer :: Average execution time after 1000 runs = 514 ns
Dot Product Index :: Average execution time after 1000 runs = 587 ns
anthony@compOrg:~/THT3$ ./performanceTestOP2
Array Size = 16
Dot Product Pointer :: Average execution time after 1000 runs = 543 ns
Dot Product Index :: Average execution time after 1000 runs = 577 ns
anthony@compOrg:~/THT3$ ./performanceTestOP3
Array Size = 16
Dot Product Pointer :: Average execution time after 1000 runs = 512 ns
Dot Product Index :: Average execution time after 1000 runs = 550 ns

```

Figure 14c: Average Execution Time Measurements for Array Size = 16

Comparing Assembly Source Files

To better explain the decrease in execution times as optimization levels increase, we must look into the compiler generated assembler source files. Figure 14d shows the compiler assembler source both for when no optimization and level 1 optimization are applied for the dot product program utilizing array indexing. The first observation here is that the assembly generated via level 1 optimization is much smaller than the assembly with no optimization. This is done by reducing the assembly required to perform an instruction (which may be more difficult to debug).

<pre> 1 .file "DotProduct_Index.cpp" 2 .text 3 .globl _Z16DotProduct_IndexPls_l 4 .type _Z16DotProduct_IndexPls_l, @function 5 _Z16DotProduct_IndexPls_l: 6 .LFB0: 7 .cfi_startproc 8 pushq %rbp 9 .cfi_def_cfa_offset 16 10 .cfi_offset 6, -16 11 movq %rsp, %rbp 12 .cfi_def_cfa_register 6 13 movq %rdi, -24(%rbp) 14 movq %rsi, -32(%rbp) 15 movl %edx, -36(%rbp) 16 movl \$0, -8(%rbp) 17 movl \$0, -4(%rbp) 18 .L3: 19 movl -4(%rbp), %eax 20 cmpl -36(%rbp), %eax 21 jge .L2 22 movl -4(%rbp), %eax 23 cltq 24 leaq 0(%rax,4), %rdx 25 movq -24(%rbp), %rax 26 addq %rdx, %rax 27 movl (%rax), %edx 28 movl -4(%rbp), %eax 29 cltq 30 leaq 0(%rax,4), %rcx 31 movq -32(%rbp), %rax 32 addq %rcx, %rax 33 movl (%rax), %eax 34 imull %edx, %eax 35 addl %eax, -8(%rbp) 36 addl \$1, -4(%rbp) 37 jmp .L3 38 .L2: 39 movl -8(%rbp), %eax 40 popq %rbp 41 .cfi_def_cfa 7, 8 42 ret 43 .cfi_endproc 44 .LFE0: 45 .size _Z16DotProduct_IndexPls_l, .-_Z16DotProduct_IndexPls_l 46 .ident "GCC: (Ubuntu 7.5.0-3ubuntu1-18.04) 7.5.0" </pre>	<pre> 1 .file "DotProduct_Index.cpp" 2 .text 3 .globl _Z16DotProduct_IndexPls_l 4 .type _Z16DotProduct_IndexPls_l, @function 5 _Z16DotProduct_IndexPls_l: 6 .LFB0: 7 .cfi_startproc 8 testl %edx, %edx 9 jle .L4 10 leal -1(%rdx), %eax 11 leaq 4(%rax,4), %r8 12 movl \$0, %edx 13 movl \$0, %eax 14 .L3: 15 movl (%rdi,%rdx), %ecx 16 imull (%rsi,%rdx), %ecx 17 addl %ecx, %eax 18 addq \$4, %rdx 19 cmpl %r8, %rdx 20 jne .L3 21 rep ret 22 .L4: 23 movl \$0, %eax 24 ret 25 .cfi_endproc 26 .LFE0: 27 .size _Z16DotProduct_IndexPls_l, .-_Z16DotProduct_IndexPls_l 28 .ident "GCC: (Ubuntu 7.5.0-3ubuntu1-18.04) 7.5.0" 29 .section .note.GNU-stack,"",@progbits </pre>
<p>-O0 (No Optimization)</p>	<p>-O1 (Optimization Level 1)</p>

Figure 14d: Comparing Assembly Source

In a similar sense, figure 14e shows the compiler generated assembly code for when optimization level 2 is applied. There are very minor differences between this code and the assembly source generated in the previous figure with level one optimizations. Most noticeably is use of `xorl %eax, %eax` instruction instead of `movl $0, %eax`. The use of `xorl` provides some optimization by requiring less space to encode as a machine instruction as it does not require the use of an immediate operand whereas the instruction via `movl` does and thus requires more space (4 Bytes).

```

1      .file      "DotProduct_Index.cpp"
2      .text
3      .p2align  4,,15
4      .globl    _Z16DotProduct_IndexPiS_i
5      .type     _Z16DotProduct_IndexPiS_i, @function
6 _Z16DotProduct_IndexPiS_i:
7 .LFB0:
8      .cfi_startproc
9      testl    %edx, %edx
10     jle      .L4
11     leal     -1(%rdx), %eax
12     xorl     %edx, %edx
13     leaq     4(,%rax,4), %r8
14     xorl     %eax, %eax
15     .p2align 4,,10
16     .p2align 3
17 .L3:
18     movl     (%rdi,%rdx), %ecx
19     imull    (%rsi,%rdx), %ecx
20     addq     $4, %rdx
21     addl     %ecx, %eax
22     cmpq     %rdx, %r8
23     jne      .L3
24     rep ret
25     .p2align 4,,10
26     .p2align 3
27 .L4:
28     xorl     %eax, %eax
29     ret
30     .cfi_endproc
31 .LFE0:
32     .size    _Z16DotProduct_IndexPiS_i, .-_Z16DotProduct_IndexPiS_i
33     .ident   "GCC: (Ubuntu 7.5.0-3ubuntu1~18.04) 7.5.0"
34     .section .note.GNU-stack,"",@progbits

```

Figure 14e: Compiler Generated Assembly Source at optimization level 2

Lastly, the application of level 3 optimization attempts to vectorize loops in an attempt to reduce execution time. However, not only does the assembly become more difficult to debug, but this application also yields the longest compilation time which may be an issue for larger programs. In addition, several sources warn against this optimization because vectorizing loops typically produces larger executables which might actually perform worse than lower optimization levels. Figure 14f shows the assembly source for level 3 optimization. Note the utilization of SSE instruction set.

```

1      .file "DotProduct_Index.cpp"
2      .text
3      .p2align 4,,15
4      .globl _Z16DotProduct_IndexPIS_i
5      .type _Z16DotProduct_IndexPIS_i, @function
6      _Z16DotProduct_IndexPIS_i:
7      .LFB0:
8      .cfi_startproc
9      testl %edx, %edx
10     jle .L9
11     movq %rdi, %r9
12     leal -1(%rdx), %ecx
13     movl $5, %r8d
14     shrq $2, %r9
15     pushq %r12
16     .cfi_def_cfa_offset 16
17     .cfi_offset 12, -16
18     pushq %rbp
19     .cfi_def_cfa_offset 24
20     .cfi_offset 6, -24
21     negq %r9
22     pushq %rbx
23     .cfi_def_cfa_offset 32
24     .cfi_offset 3, -32
25     andl $3, %r9d
26     leal 3(%r9), %eax
27     cmpl $5, %eax
28     cmovb %r8d, %eax
29     cmpl %eax, %ecx
30     jb .L10
31     testl %r9d, %r9d
32     je .L11
33     movl (%rdi), %eax
34     movl $1, %ebp
35     inull (%r12), %eax
36     cmpl $1, %r9d
37     movl %eax, %r12d
38     je .L4
39     movl 4(%rdi), %eax
40     movl $2, %ebp
41     inull 4(%r12), %eax
42     addl %eax, %r12d
43     cmpl $2, %r9d
44     je .L4
45     movl 8(%r12), %eax
46     movl $3, %ebp
47     inull 8(%rdi), %eax
48     addl %eax, %r12d
49     .L4:
50     movl %edx, %ebx
51     pxor %xmm3, %xmm3
52     subl %r9d, %ebx
53     salq $2, %r9
54     xorl %ecx, %ecx
55     movl %ebx, %r11d
56     leaq (%rdi,%r9), %r10
57     xorl %r8d, %r8d
58     shrq $2, %r11d
59     addq %r12, %r9
60     .p2align 4,,10
61     .p2align 3
62     .L7:
63     movdqu (%r9,%rcx), %xmm0
64     addl $1, %r8d
65     movdqa %xmm0, %xmm1
66     psrlq $32, %xmm0
67     movdqa (%r10,%rcx), %xmm2
68     pmuludq (%r10,%rcx), %xmm1
69     pshufd $8, %xmm1, %xmm1
70     addq $16, %rcx
71     psrlq $32, %xmm2
72     pmuludq %xmm2, %xmm0
73     pshufd $8, %xmm0, %xmm0
74     cmpl %r11d, %r8d
75     punpckldq %xmm0, %xmm1
76     paddb %xmm1, %xmm3
77     jb .L7
78     movdqa %xmm3, %xmm0
79     movl %ebx, %r8d
80     andl $-4, %r8d
81     psrlq $8, %xmm0
82     paddb %xmm0, %xmm3
83     movdqa %xmm3, %xmm0
84     leal (%r8,%rbp), %ecx
85     psrlq $4, %xmm0
86     paddb %xmm0, %xmm3
87     movd %xmm3, %eax
88     addl %r12d, %eax
89     cmpl %r8d, %ebx
90     je .L1
91     .L3:
92     movslq %ecx, %r9
93     movl (%rdi,%r9,4), %r8d
94     inull (%r12,%r9,4), %r8d
95     addl %r8d, %eax
96     leal 1(%rcx), %r8d
97     cmpl %r8d, %edx
98     jle .L1
99     movslq %r8d, %r8
100    movl (%rdi,%r8,4), %r9d
101    inull (%r12,%r8,4), %r9d
102    leal 2(%rcx), %r8d
103    addl %r9d, %eax
104    cmpl %r8d, %edx
105    jle .L1
106    movslq %r8d, %r8
107    movl (%rdi,%r8,4), %r9d
108    inull (%r12,%r8,4), %r9d
109    leal 3(%rcx), %r8d
110    addl %r9d, %eax
111    cmpl %r8d, %edx
112    jle .L1
113    movslq %r8d, %r8
114    movl (%rdi,%r8,4), %r9d
115    inull (%r12,%r8,4), %r9d
116    leal 4(%rcx), %r8d
117    addl %r9d, %eax
118    cmpl %r8d, %edx
119    jle .L1
120    movslq %r8d, %r8
121    addl $5, %ecx
122    movl (%rdi,%r8,4), %r9d
123    inull (%r12,%r8,4), %r9d
124    addl %r9d, %eax
125    cmpl %ecx, %edx
126    jle .L1
127    movslq %ecx, %rcx
128    movl (%rdi,%rcx,4), %edx
129    inull (%r12,%rcx,4), %edx
130    addl %edx, %eax
131    .L1:
132    popq %rbx
133    .cfi_restore_state
134    .cfi_def_cfa_offset 24
135    popq %rbp
136    .cfi_def_cfa_offset 16
137    popq %r12
138    .cfi_def_cfa_offset 8
139    ret
140    .p2align 4,,10
141    .p2align 3
142    .L11:
143    .cfi_restore_state
144    xorl %ebp, %ebp
145    xorl %r12d, %r12d
146    jmp .L4
147    .p2align 4,,10
148    .p2align 3
149    .L9:
150    .cfi_def_cfa_offset 8
151    .cfi_restore 3
152    .cfi_restore 6
153    .cfi_restore 12
154    xorl %eax, %eax
155    ret
156    .p2align 4,,10
157    .p2align 3
158    .L10:
159    .cfi_def_cfa_offset 32
160    .cfi_offset 3, -32
161    .cfi_offset 6, -24
162    .cfi_offset 12, -16
163    xorl %ecx, %ecx
164    xorl %eax, %eax
165    jmp .L3
166    .cfi_endproc
167    .LFE0:
168    .size _Z16DotProduct_IndexPIS_i, .-_Z16DotProduct_IndexPIS_i
169    .ident "GCC: (Ubuntu 7.5.0-3ubuntu1-18.04) 7.5.0"
170    .section .note.GNU-stack,"",@progbits

```

Figure 14f: Compiler Generated Assembly Source at optimization level 3

Results

Similar observations occurred when examining the program utilizing pointer arithmetic. Figures 51a and 15b show the plotted execution times utilizing various optimization levels for various array sizes.

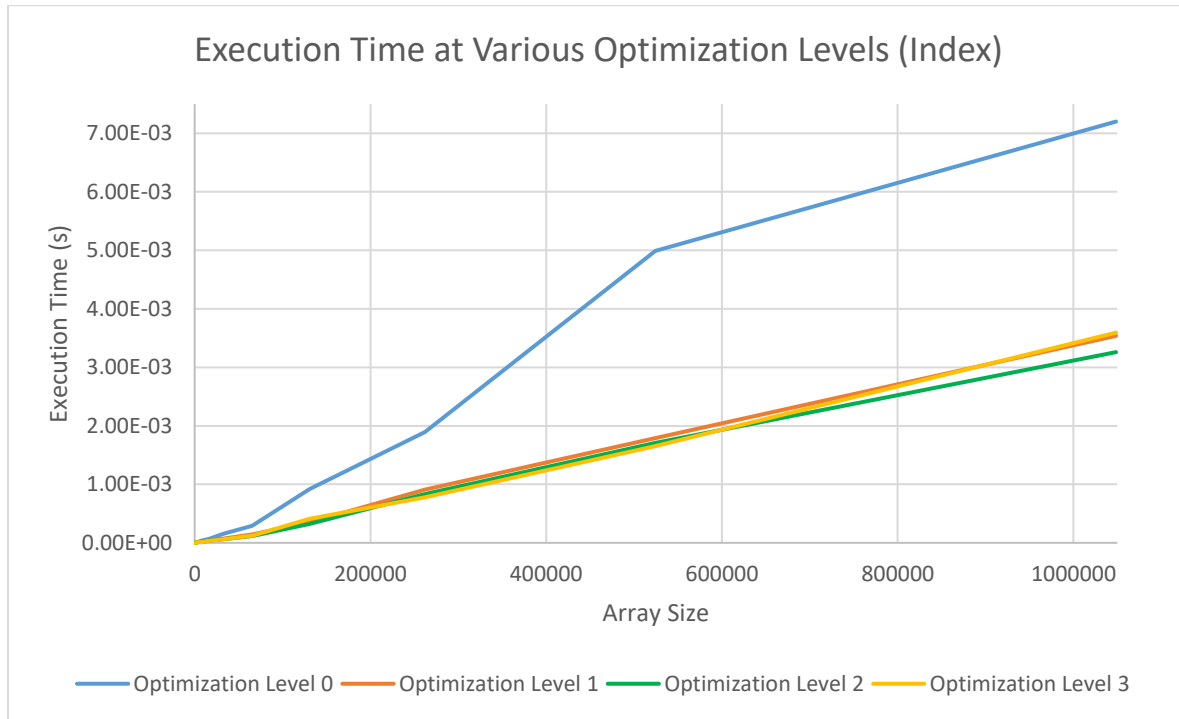


Figure 15a: Execution Times at Various Optimization Levels (Index)




Figure 15b: Execution Times at Various Optimization Levels (Pointer)

Given these results, it is obvious that applying one level of optimization significantly improves the execution time of both programs. It can also be concluded that optimization at level 2 provides even faster execution times than level 1 optimization. However, the results obtained from optimization level 3, failed to prove any improvement over level 2 optimization in regard to execution time. In fact, there were some instances where this optimization level performed worse than level 2 optimization (but never worse than level 1). In addition to no improvement in execution time, we pay the price with higher compilation time which was noticeable for larger array sizes. This behavior confirms what online sources previously stated. Hence it could be concluded that level 2 optimization is the much better option in terms of reducing execution time (if a slight increase in compilation time is no issue). Regarding Level 1 optimization, it is a good choice if maximum optimization is not required.